Computer and robot vision

Course notes

Daniele Besozzi

A.A. 26/27

Contents

1 Image formation 2
1.1 Pinhole camera model 2
1.1.1 Simplified pinhole model 3
1.1.2 General pinhole model 5
1.1.3 Intrinsic parameters 5
1.1.4 Extrinsic parameters 6
1.1.5 Thin lenses 7
1.2 Calibration 9
1.2.1 Direct linear transform method (Tsai) 9
1.2.2 Reprojection error 10
1.3 Zang's method 10
1.4 Sensing 11
1.4.1 CCD sensors 11
1.4.2 CMOS sensors 13
1.4.3 Sensor models 14
1.5 Digital images 15
1.5.1 HDR 15
1.5.1.1 Multiple exposure 16
1.5.1.2 Single exposure 16
1.5.2 Color cameras 17
1.5.2.1 Color filter array 17
1.5.2.2 3-CCD cameras 18
1.5.2.3 Stacked sensors (Foveon) 18
1.6 3D cameras 18
1.6.1 Time-of-flight cameras 19
1.7 Defects in images 19
1.7.1 Defects caused by the sensor 19
1.7.1.1 Blooming 19
1.7.1.2 Smear 19
1.7.2 Defects caused by the lens 20
1.7.2.1 Chromatic aberration 21
1.7.2.2 Vignetting 22
B. Mathematical concepts 23
B.1. Roto-translation 23
B.2. Euler angles 24
B.3. Roll, pitch and yaw angles 24
B.4. Roto-translation matrix 25

1 Image formation

A camera is the union of the sensors and the software that uses the image to determine relevant points in the image plane and the projection model. This allows to calculate the lines connecting the points in the image plane and the object.

> A camera sensor (once calibrated) does not measure points in the image, but measures lines in the space (a subset of the star of lines passing trough the pinhole).

Since we are only calculating lines, knowing the precise physical measures of the image plane is not necessary since the lines don't depend on them. The simplest geometric model of image formation is the **pinhole camera**, it is based on the same principle as the camera obscura.

1.1 Pinhole camera model

The tiny hole (pinhole) in one wall of the room lets in a ray of light for each point in the scene so that an image of the outside world is drawn on the opposite wall (**image plane**).

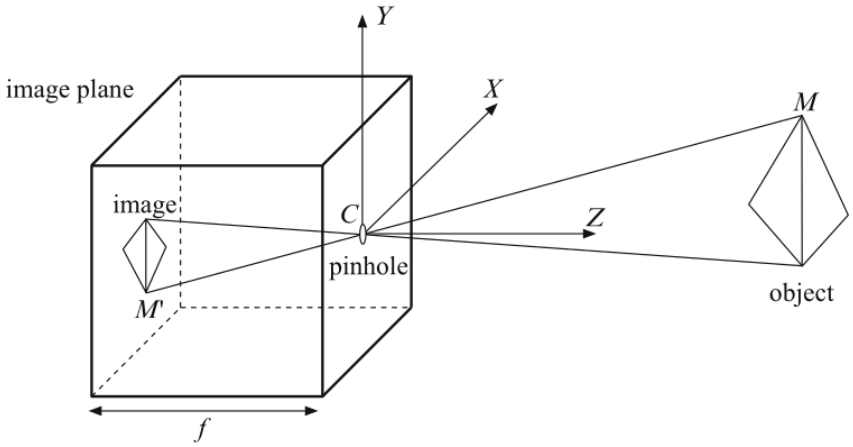


Figure 1: Image formation in the camera obscura/pinhole camera

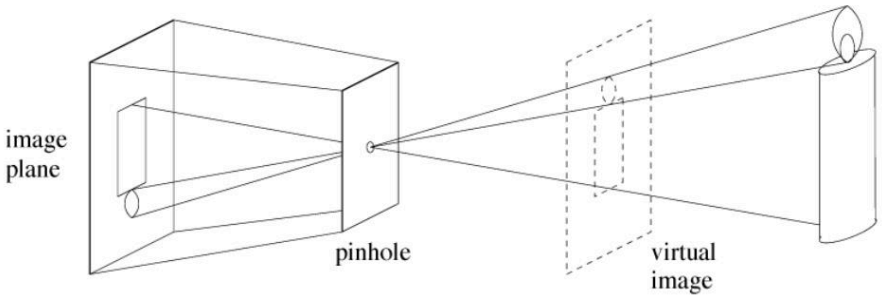


Figure 2: Pinhole camera with virtual image

The geometric model of the camera that we will adopt is the so-called pinhole or perspective model, which consists (refer to Figure 3) of the **image plane** Q and the **centre of Projection (COP)** C, located at a distance f (**focal length**) from the image plane.

> **Definition 1.1** *(principal axis).* The line through C orthogonal to Q is the **principal axis** of the camera

> **Definition 1.2** *(principal point).* The intersection of the principal axis with Q is called the **principal point** (or center of the image).

Definition 1.3 *(focal plane).* The plane F parallel to Q and containing C is called the **focal plane**.

The points in F are projected to infinity on Q.

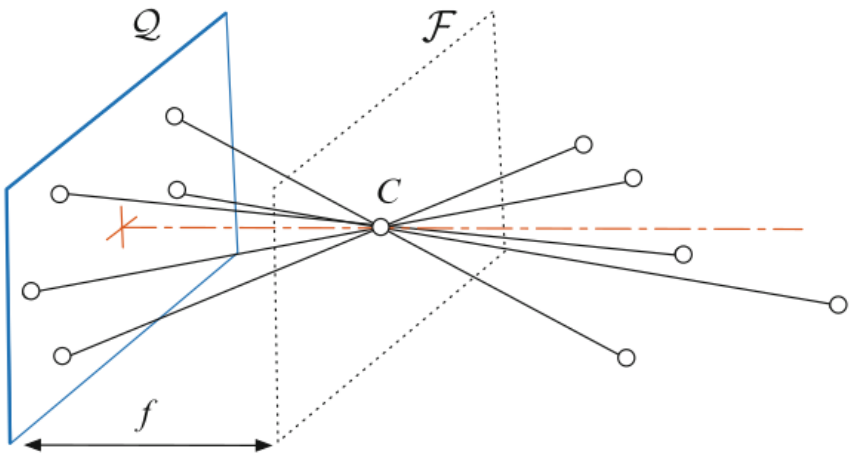


Figure 3: Pinhole camera with virtual image

Let M be a point in the scene, of coordinates (X, Y, Z), and let M' of coordinates (X', Y', Z') be its projection onto the image plane through the pinhole C. If f is the distance of C from the image plane (focal length), then from the similitude of the triangles we get:

$$\frac{-X'}{f} = \frac{X}{Z} \qquad \frac{-Y'}{f} = \frac{Y}{Z} \tag{1}$$

therefore

$$X' = \frac{-fX}{Z} \qquad Y' = \frac{-fY}{Z} \qquad Z' = -f \tag{2}$$

Note that the image is inverted with respect to the scene, both left-right and top-bottom, as indicated by the minus sign. These equations define the image formation process, which is called **perspective projection**.

The parameter f determines the **magnification factor** of the image: if the image plane is close to the pinhole, the image is reduced in size, while it will become larger and larger as f increases. Since the working area of the image plane is framed, moving the image plane away also reduces the **field of view**, that is the solid angle that contains the portion of the scene that appears in the image.

The division by Z is responsible for the foreshortening effect, whereby the apparent size of an object in the image decreases according to its distance from the observer.

1.1.1 Simplified pinhole model

The coordinate system in which the points of the three-dimensional space are represented is called the **world (or object) coordinate system**. Let us introduce the **camera reference frame**, a right-handed triplet (X, Y, Z) of axes centred in C and with the Z axis coincident with the principal axis. In this simplified model, the world reference frame is chosen coincident with the camera reference frame.

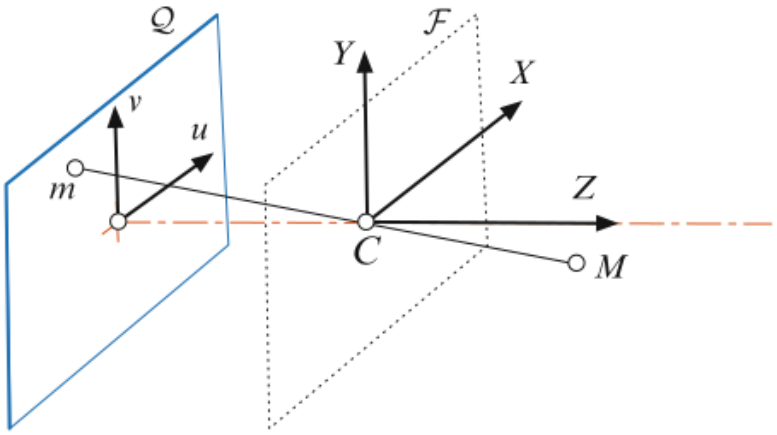


Figure 4: Simplified model of the camera with reference frames. The point M is projected onto the point m where the line CM intersects the image plane Q

We also introduce a **reference frame** (u, v) for the plane Q centred at the principal point and with the axes u and v oriented as X and Y , respectively, as shown in Figure 4.

Unless stated otherwise, we use ~ to denote Cartesian coordinates to distinguish them from homogeneous ones.

Now consider a point M of coordinates $\widetilde{M} = \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$ in 3D space and let m of coordinates $\tilde{m} = \begin{pmatrix} u \\ v \end{pmatrix}$ be its projection onto Q trough C. We have that

$$\begin{cases} u = \frac{-f}{Z} X \\ v = \frac{-f}{Z} Y \end{cases} \tag{3}$$

In Equation 3 we have the intersection of two planes, therefore we obtain a line. This is the **perspective projection**. The transformation in Cartesian coordinates is clearly non-linear because of the division by Z. But if we use homogeneous coordinates, the transformation becomes linear. Let

$$m = \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} \quad M = \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} \tag{4}$$

be the homogeneous coordinates of m and M, respectively. So the equation for perspective projection, in this simplified case, rewrites:

$$Z \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = Z \begin{pmatrix} -f\frac{X}{Z} \\ -f\frac{Y}{Z} \\ Z \end{pmatrix} = \underbrace{\begin{pmatrix} -f & 0 & 0 & 0 \\ 0 & -f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix}}_{P} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} \tag{5}$$

In other words

$$m = \frac{1}{Z} P M \tag{6}$$

The matrix P represents the geometric model of the camera and is called **Perspective Projection Matrix (PPM)** or **camera matrix**.

1.1.2 General pinhole model

A realistic model of a camera, describing the transformation from 3D coordinates to pixel coordinates, as well as the perspective transformation, must take into account:

- The change of coordinates from metres in the sensor plane to pixels in the image.
- The change of coordinates between the camera reference frame and the world reference frame.

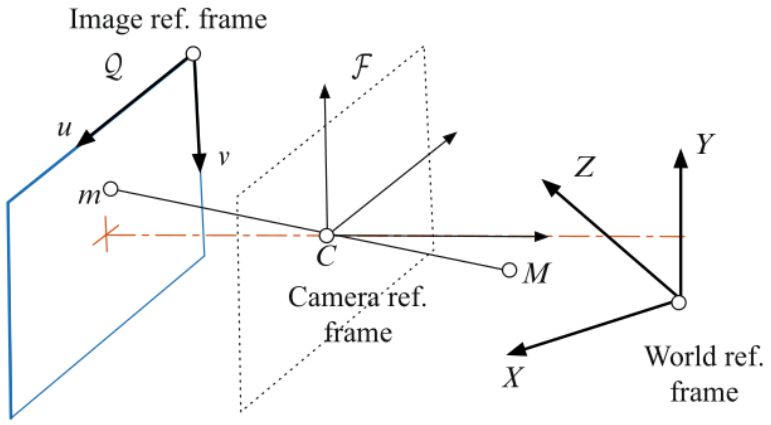
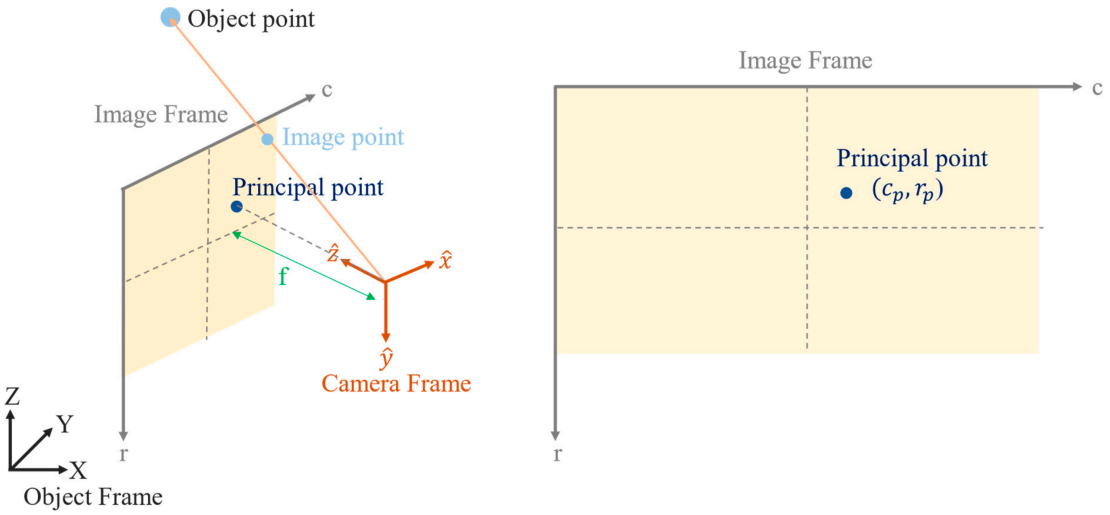


Figure 5: General model of the camera with reference frames. Note that the image plane should be viewed from the half-space opposite to C, so that the origin of the image reference is in the upper left corner of the image



(a) (b)

Figure 6: Geometry of perspective projection: (a) the three coordinate frames involved and (b) the image coordinates of the principal point.

1.1.3 Intrinsic parameters

The transformation from the sensor coordinates to pixels is modelled by an affine transformation A that takes into account the **translation of the origin at the upper left corner**, the independent rescaling of the u and v axes for changing the unit of measurement from metres to pixels and the reversal of the direction of the axes:

$$A = \begin{pmatrix} -k_u & 0 & u_0 \\ 0 & -k_v & v_0 \\ 0 & 0 & 1 \end{pmatrix} \tag{7}$$

where (u_0, v_0) are the coordinates in pixels of the principal point (which is not on the vertical axis!) and $k_u(k_v)$ (which is translated in respect to the center of the lens) is the inverse of the effective pixel size along the direction of the u (v) axis, measured in m^{-1}.

The matrix A transforms the coordinates from metres to pixels; hence, it pre-multiplies the PPM (refer to Section B.1.), which becomes:

$$P = \begin{pmatrix} -k_u & 0 & u_0 \\ 0 & -k_v & v_0 \\ 0 & 0 & 1 \end{pmatrix} \underbrace{\begin{pmatrix} -f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix}}_{\text{PPM}} = \begin{pmatrix} fk_u & 0 & u_0 & 0 \\ 0 & fk_v & v_0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix} = K(I_3 \ \ 0) \tag{8}$$

where

$$K = \begin{pmatrix} fk_u & 0 & u_0 \\ 0 & fk_v & v_0 \\ 0 & 0 & 1 \end{pmatrix} \tag{9}$$

is called the **intrinsic parameters matrix** of the camera, while $(I_3 \ \ 0)$ is the perspective transformation without parameters. It is obtained when the world reference frame coincides with the camera reference frame and the image points are expressed in normalized image coordinates, that is,

$$q = K^{-1}m \tag{10}$$

> Note that whereas pixel coordinates are always measurable in the image, to step to normalized image coordinates the matrix of intrinsic parameters is needed.

> The multiplication by the PPM matrix calculates the projection of the point in the image plane, while the multiplication by matrix A does the translation with respect to the principal point. The multiplication by the PPM matrix returns a 3×1 matrix with the coordinates of the point in the image plane expressed in metres, while the multiplication by matrix A returns a 3×1 matrix with the coordinates of the point in pixels.

A more general K includes also the **skew** angle Θ that corresponds to the angle between the u and v axes, that normally is $\frac{\pi}{2}$:

$$K = \begin{pmatrix} fk_u & \frac{-fk_u}{\tan(\Theta)} & u_0 \\ 0 & \frac{fk_v}{\sin(\Theta)} & v_0 \\ 0 & 0 & 1 \end{pmatrix} \tag{11}$$

This matrix has five degrees of freedom, which are the intrinsic parameters of the camera: fk_u, fk_v, u_0, v_0, and Θ.

> The skew angle Θ is usually very close to $\frac{\pi}{2}$, so it is often neglected in practice.

1.1.4 Extrinsic parameters

To account for the fact that—in general—the world reference frame does not coincide with the camera reference frame, we must introduce a coordinate change consisting of a direct **isometry**, that is, a **rotation** $\boldsymbol{R}$ and a **translation** $\boldsymbol{t}$. We denote by M_c the homogeneous coordinates of a point in the camera coordinate system and M the homogeneous coordinates of the same point in the world coordinate system. We can therefore write:

$$M_c = GM \tag{12}$$

where

$$G = \begin{pmatrix} R & t \\ 0 & 1 \end{pmatrix} \tag{13}$$

The matrix G encodes the **exterior orientation** (position and angular attitude[1]) of the camera with respect to the world reference frame. Specifically, the rows of R are the X, Y and Z axes of the camera reference frame expressed in the world coordinate system.

By appropriately parameterising the rotation R, the exterior orientation of the camera is encoded by six **extrinsic parameters**:

- three Euler angles ω, φ, κ that specify the rotation
- three components t_1, t_2, t_3 for the translation.

From Equation 8 and Equation 11 we can write the general perspective projection matrix as:

$$P = K[I_3, 0]G \tag{14}$$

This is the most general form of the perspective projection matrix:

- G encodes the exterior orientation of the camera
- K contains its intrinsic parameters (by analogy, we say that it encodes the interior orientation)
- $[I_3, 0]$ represents a perspective transformation from camera coordinates to normalised image coordinates.

In total, the perspective projection matrix has **11 degrees of freedom**:

- 6 for the exterior orientation
- 5 for the interior orientation.

1.1.5 Thin lenses

A pinhole camera is really small (in theory, infinitesimal) therefore the energy reaching the image plane also is. In order to obtain an image we would need a transformation rate which is, in theory, infinite and therefore impossible to obtain. If we simply increased the size in the hole we would not be able to control the amount of energy coming to the image plane, and would therefore obtain a spot instead of a point the image plane. If this were to occur, then the image would be blurry (**circle of confusion**).

Vertebrate eyes and cameras use lenses. A lens, being larger than a pinhole, can gather more light. Lenses allow a larger opening, so a finite amount of energy, to reach the image plane and at the same time we might be able to obtain an effect in terms of projection (zoom, etc...). The downside is that not the whole scene can be in focus at the same time.

The approximation we make for the optics of the imaging system— which in general is very complex, being made up of multiple lenses is that of the thin lens.

> **Definition 1.4** *(Thin lenses).* Thin lenses are ideal devices that enjoy the following properties:
>
> - The rays parallel to the optical axis incident on the lens are refracted so that they pass through a point on the optical axis called focus F
> - The rays passing through the centre C of the lens are unaffected.

The distance of the focus F from the centre of the lens C is called **focal length** D, which depends on the radii of curvature of the two lens surfaces and the refractive index of the constituent material.

[1] In navigation orientation is the determination of the east cardinal point (oriens in latin); by extension, in topography and photogrammetry, it indicates the determination of position and direction with respect to cardinal points or other reference frames.

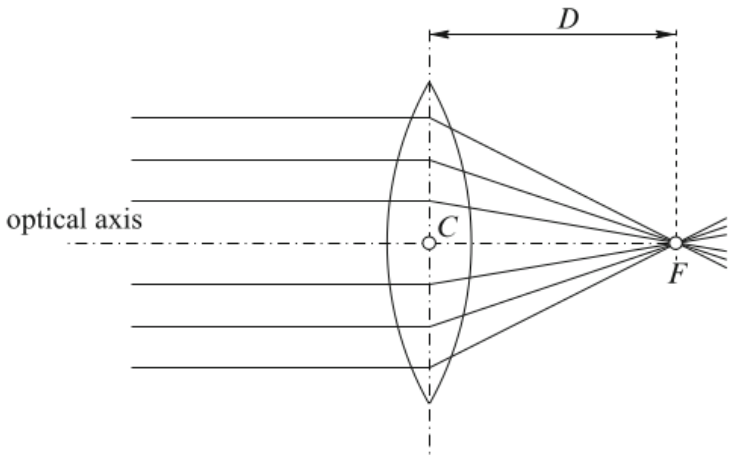


Figure 7: Thin lenses

> For non-thin lenses there is no "center" of the lens, but two nodal points (front and rear) are defined to characterize its behavior. By definition, an input ray directed at the front nodal point leads to an output ray which appears to come from the back nodal point, forming the same angle with respect to the optical axis. For a thin lens, the two nodal points coincide in the centre of the lens.

Given a point of the scene M it is possible to construct graphically the image (or conjugate point) M', using two special rays that depart from M: the ray parallel to the optical axis, which after refraction, passes through F and the ray that passes unchanged through C.

Due to this construction and the similitude of the triangles, we obtain the conjugate points formula (or equation of the thin lens):

$$\frac{1}{Z} + \frac{1}{Z'} = \frac{1}{D} \tag{15}$$

The image of a point in the scene at distance Z from the lens is produced (in focus) at a distance Z', from the lens, which depends on the depth Z of the point and the focal length D of the lens.

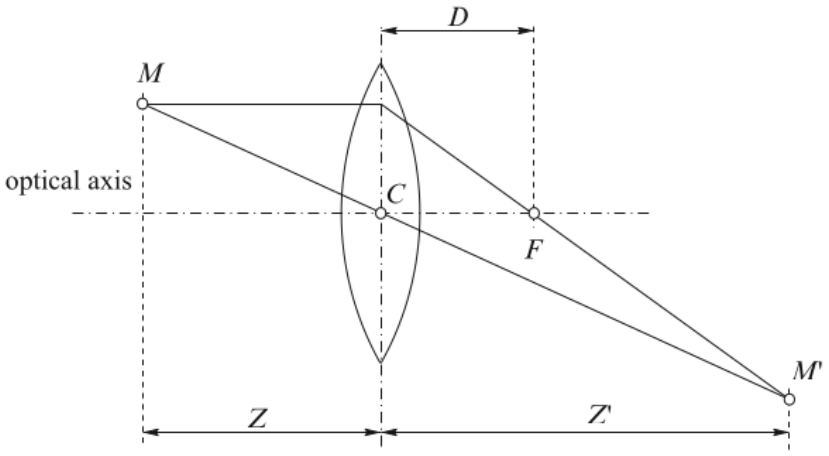


Figure 8: Construction of the conjugate point.

The image of the point M, when in focus, is formed on the image plane at the same point predicted by the pinhole model with the pinhole coincident with the centre of the lens C: in fact, the ray passing through C is common to the two geometric constructions. The other light rays that leave M and are collected by the lens to increase the light reaching M'.

If Equation 15 is not verified one gets a blurred image of the point, that is, a circle which is named **circle of confusion**.

> This is the same effect that would occur by just increasing the size of the pin hole, but is caused by different physical conditions.

The photosensitive elements have a small but finite size.

Remark 1.5. As long as the circle of confusion does not exceed the size of the photosensitive element, the image is in focus. This is due to the **sampling theorem**.

Therefore, there is an interval of depths for which the points are in focus. This range is called **depth of field**, which is inversely proportional to the diameter of the lens; in fact, the pinhole camera has an infinite depth of field. The light that is collected, on the other hand, is directly proportional to the diameter of the lens.

Remark 1.6. The focal length of the lens is a different notion from the focal length of the pinhole camera. Two different notations, f and D, were used on purpose, although, unfortunately, the two quantities have the same name.

1.2 Calibration

We now address the problem of precisely and accurately determining the parameters of the pinhole model (intrinsic and extrinsic), a process that goes by the name of **calibration of the camera**. The idea of the resection methods for calibration is that by knowing the image projections of 3D points of known coordinates (**control points**), it should be possible to obtain the unknown parameters by solving the perspective projection equation. K, R, T can ben obtained by applying the perspective projection equation to the control points and their projections in the image.

1.2.1 Direct linear transform method (Tsai)

Given the correspondences between n non-coplanar control points M_i and their projections in the image m_i, it is required to compute a matrix P such that:

$$m_i \approx PM_i \quad i = 1, ..., n \tag{16}$$

The image coordinates of the control points can be obtained manually or by classical image processing methods. To eliminate the scaling factor, we exploit the cross product, rewriting the previous equation as:

$$m_i \times PM_i = 0 \tag{17}$$

Skipping some algebraic manipulations, we can rewrite the previous equation as a linear system of equations in the unknowns $p_1, ..., p_{\{12\}}$, which are the entries of the matrix P. Only 2 of the three equations are linearly independent. With n points we obtain a system of $2n$ homogeneous linear equations, which we rewrite as

$$A \ \mathsf{vec} \ p = 0 \tag{18}$$

where p is the vector of unknowns and A is a $2n \times 12$ matrix of coefficients, which depend on the coordinates of the control points.

> In theory, then, six non-coplanar points are sufficient for the computation of P ; if any four of them are coplanar, the rank of A drops and the DLT algorithm fails.
>
> In practice, it is advisable to use all the available points (the more, the better) to compensate for the inevitable measurement errors.

We solve Equation 18 with the least-squares method.

1.2.2 Reprojection error

THe reprojection error is the Euclidean distance (in pixels) between an observed image point and the corresponding 3D point reprojected onto the camera frame. This give us a quantitative measure of the accuracy of the calibration (ideally 0).

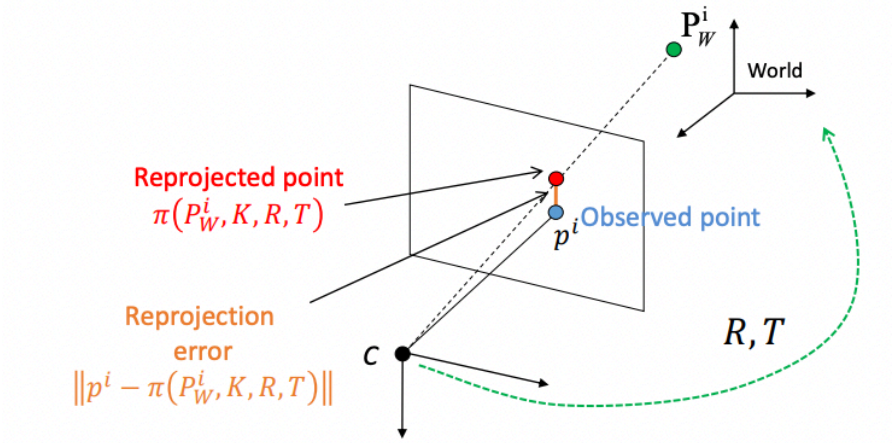


Figure 9: Reprojection of a point.

We can optimize the calibration by minimizing the reprojection error, which is a non-linear optimization problem, in which we also include the lens distortion parameters. This is usually done using an iterative method such as the Levenberg-Marquardt algorithm.

1.3 Zang's method

This method is the state of the art of camera calibration, and is the most used in practice. It is based on the use of **multiple views of a planar calibration pattern**, which is easier to manufacture and use than a 3D calibration object.

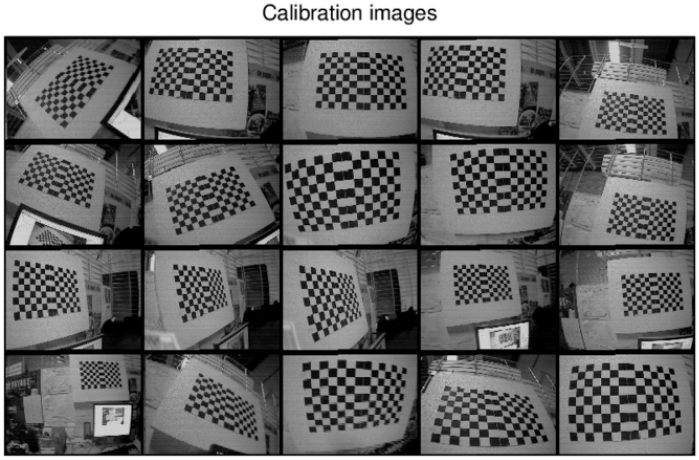


Figure 10: Zhang's method for camera calibration. The calibration pattern is a planar checkerboard, which is easy to manufacture and use.

First, we need to compute where are the corners of the checkerboard (purely image processing problem). This is done because there are many pixels on the corners of the boxed, by interpolating the corners we smooth out the noise.

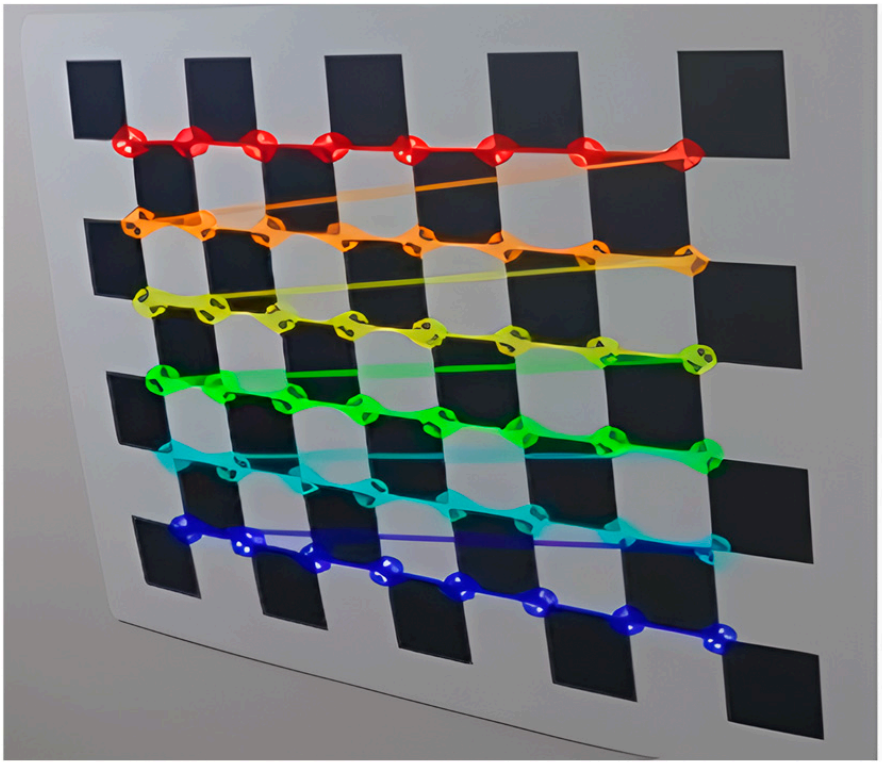


Figure 11: Zhang's method for camera calibration. The corners of the checkerboard are detected in the image.

Then we need to compute the homography between the checkerboard and the image plane, which is a linear problem.

1.4 Sensing

What differentiates a camera (in the modern sense of the word) from the portable camera obscura of the 17th century is its ability to record the pictures that form on its backplane. We now introduce two different kinds of modern camera sensors.

1.4.1 CCD sensors

Charge-coupled-device (CCD) cameras that were proposed in 1970. A CCD sensor uses a rectangular grid of electron-collection sites laid over a thin silicon wafer to record a measure of the amount of light energy reaching each of them.

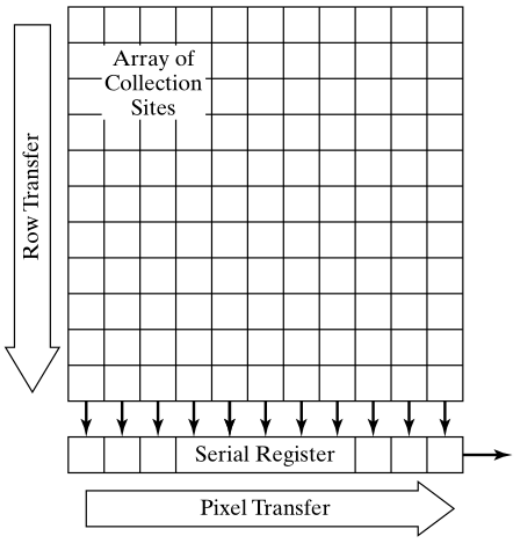


Figure 12: A CCD Device.

Each site is formed by growing a layer of silicon dioxide on the wafer and then depositing a conductive gate structure over the dioxide. When photons strike the silicon, electron-hole pairs are generated (photo-conversion), and the electron are captured by the potential well formed by applying a positive electrical potential to the corresponding gate. The electrons generated at each site are collected over a fixed period of time T.

At this point, the charges stored at the individual sites are moved using charge coupling:

- Charge packets are transferred from site to site by manipulating the gate potentials, preserving the separation of the packets.
- The image is read out of the CCD one row at a time, each row being transferred in parallel to a serial output register with one element in each column.
- Between two row reads, the register transfers its charges one at a time to an output amplifier that generates a signal proportional to the charge it receives.

This process continues until the entire image has been read out. The reading time depends on the use case, varying greatly.

> The digital output of most CCD cameras is transformed internally into an analog video signal before being passed to a frame grabber that constructs the final digital image.

Consumer-grade color CCD cameras essentially use the same chips as black-and-white cameras, except that successive rows or columns of sensors are made sensitive to red, green or blue light often using a filter coating that blocks the complementary light.

The spatial resolution of single-CCD cameras is of course limited, and higher-quality cameras use a beam splitter to ship the image to three different CCDs via color filters.

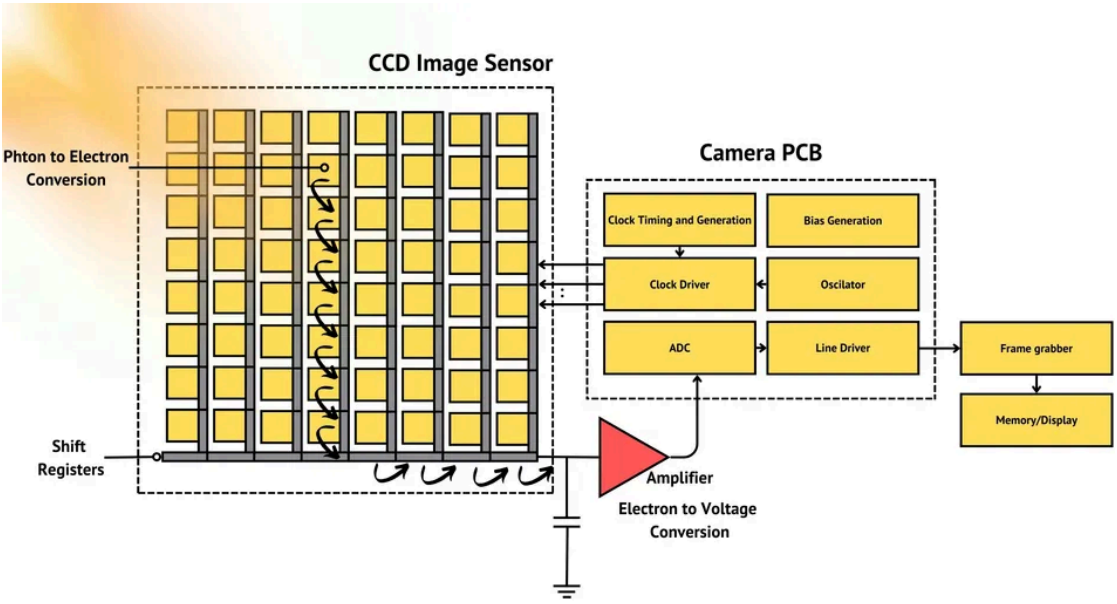


Figure 13: CCD image sensor pipeline

Remark 1.7. The flow of electricity in CCD sensors obviously generates heat, which equals energy. The sensors pick up this heat as noise, therefore the sensors need to be cooled to reduce noise.

1.4.2 CMOS sensors

CMOS sensors are integral to modern digital imaging, particularly in devices like smartphones, digital cameras, and high-speed industrial systems. Unlike CCD sensors, CMOS sensors process image data in a way that allows for faster operation and lower power consumption.

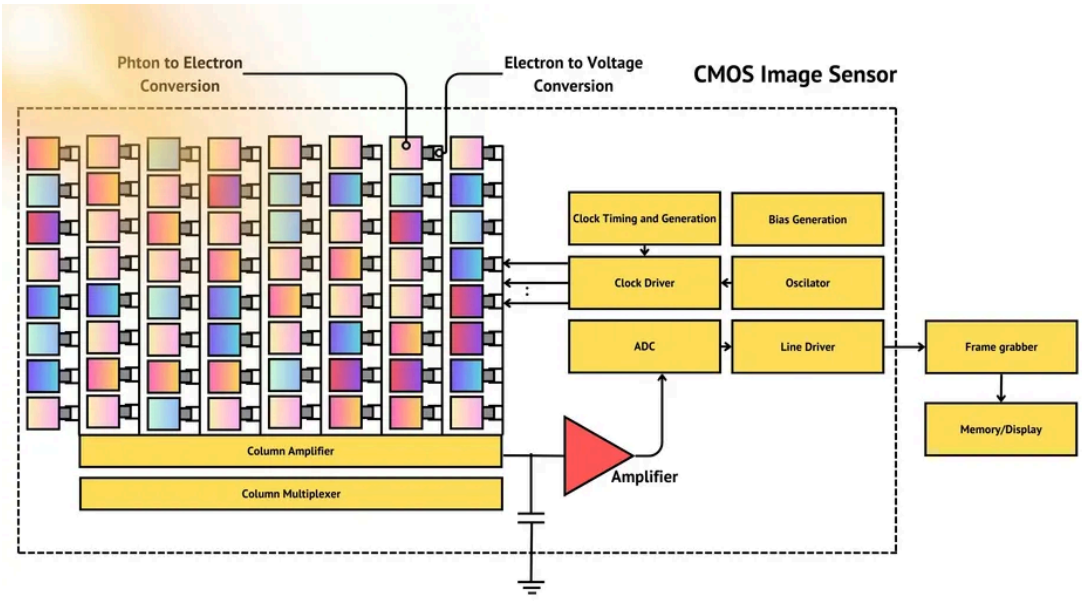


Figure 14: CMOS image sensor pipeline

Just like in a CCD sensor, the process starts when light enters the camera lens and is focused onto the CMOS sensor. The CMOS sensor is made up of pixels arranged in a grid. Each pixel acts as a light detector, capturing photons and converting them into an electrical charge. When photons hit the CMOS sensor, they are absorbed by the photodiodes within each pixel. This absorption generates electron-hole pairs in the silicon substrate, where each photon generates one electron. The number of electrons generated in each pixel is proportional to the

intensity of the light striking that pixel. Brighter areas in the image produce more electrons, corresponding to higher brightness levels.

After the charge is converted to voltage at the pixel level, the signal is amplified to make it strong enough for further processing. This amplification occurs within the same pixel, minimizing the distance the signal has to travel, which reduces noise and improves image quality. In many CMOS sensors, the analog voltage signals are converted to digital signals at the column or even the pixel level by an integrated ADC. This localized conversion is one of the reasons CMOS sensors can achieve high-speed readouts.

One characteristic of many CMOS sensors is the **rolling shutter** effect. Because each row or column of pixels is read out sequentially rather than all at once, moving objects can appear distorted. For example, a fast-moving car might look skewed because different parts of the image are captured at slightly different times. However, advanced CMOS sensors can mitigate this effect, and some even employ a **global shutter** to capture all pixels simultaneously, similar to CCD sensors.

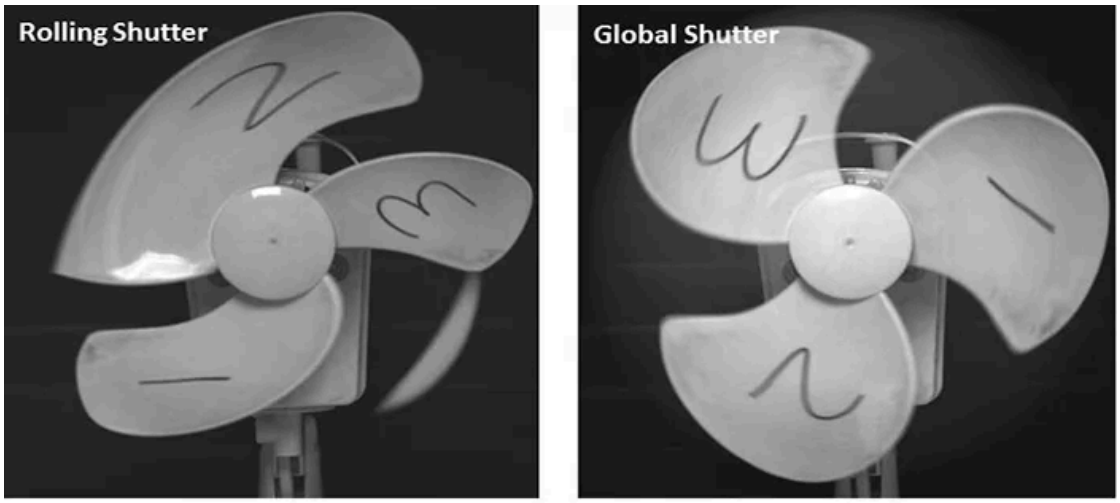


Figure 15: Rolling and global shutter example of the same image.

> Optimally we want a global shutter with an exposure time that doesn't lead to motion blur.

1.4.3 Sensor models

For simplicity, we restrict our attention in this section to black-and-white CCD cameras: Color cameras can be treated in a similar fashion by considering each color channel separately and taking the effect of the associated filter response explicitly into account.

The number I of electrons recorded at the cell located at row $r$and column c of a CCD array can be modeled as

$$I(r,c) = T \int_{\lambda} \int_{p \in S(r,c)} E(p,\lambda) R(p) q(\lambda) dp \; d\lambda \tag{19}$$

where T is the electron-collection time and the integral is calculated over the spatial domain $S(r,c)$ of the cell and the range of wavelengths to which the CCD has a nonzero response. In this integral:

- E is is the power per unit area and unit wavelength arriving at the point p
- R is the spatial response of the site
- q is the quantum efficiency of the device

Definition 1.8 *(Quantum efficiency of a device).* The quantum efficiency of a device is the number of electrons generated per unit of incident light energy

In general, both E and q depend on the light wavelength λ, and E and R depend on the point location p within $S(r, c)$.

The output amplifier of the CCD transforms the charge collected at each site into a measurable voltage. In most cameras, this voltage is then transformed into a low-pass-filtered (generally a mean) video signal by the camera electronics with a magnitude proportional to I. The analog image can be once again transformed into a digital one using a frame grabber that spatially samples the video signal and quantizes the brightness value at each image point or pixel (from picture element).

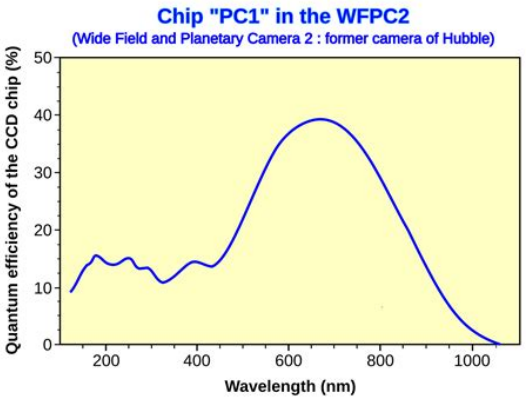


Figure 16: quantum efficiency of Hubble camera 2

In addition to this, we can create a diagram that shows us the relation between the amount of energy incoming into the device (independently of the wavelength) and the "amount" of signal coming out of the device, the so-called **in/out diagram**. Of course, we ideally want it to be linear.

1.5 Digital images

Here we introduce digital images and some operations that can be done on them.

1.5.1 HDR

According to exposure time, length of exposure time and the scene itself some areas of the image can be overexposed, underexposed, or correctly exposed. We want to obtain an image in which all areas are properly exposed. This is called **High Dynamic Range (HDR)** imaging. Of course, the approach changes whether the image is static or dynamic.

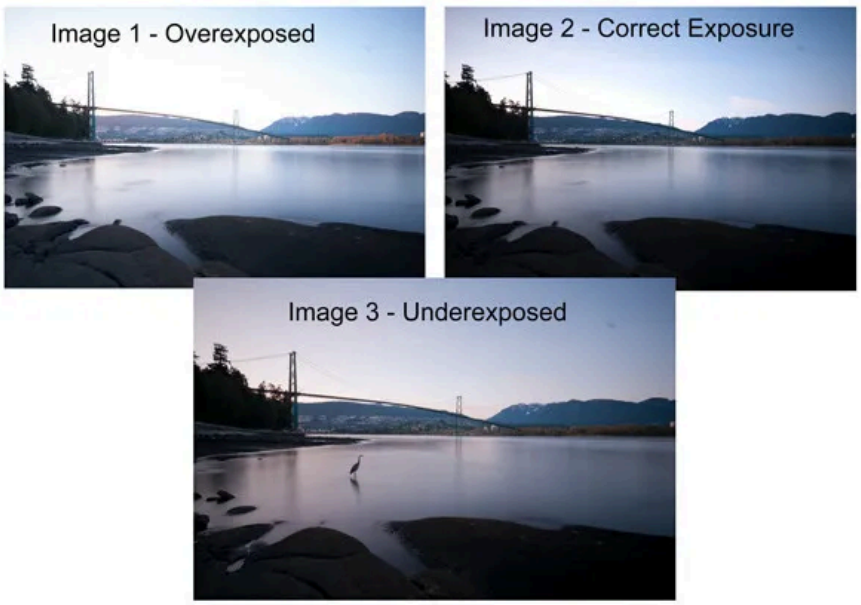


Figure 17: pictures taken with different exposure times of the same scene

There are two main approaches to HDR imaging:

1.5.1.1 Multiple exposure

In this approach, multiple very fast and near exposures are taken on the same scene, with different exposure conditions. These images are combined after being read from the sensor, the algorithm may be heavy and the image can be shown only after the processing is done. This may be acceptable depending on the use case.

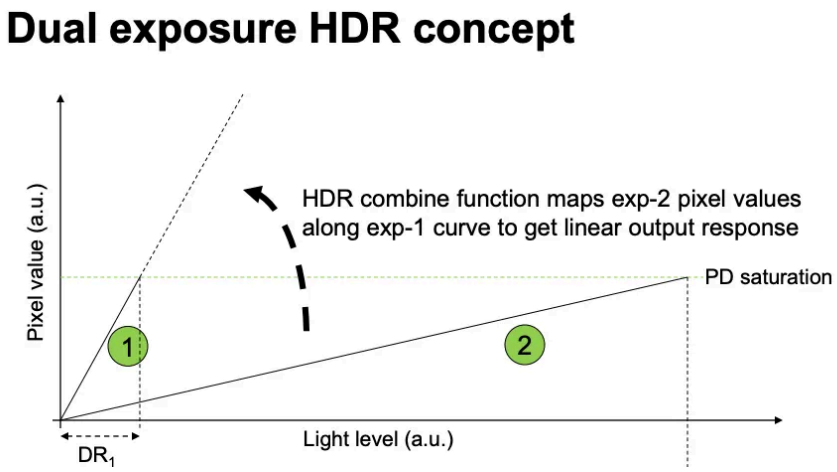


Figure 18: Multiple exposure HDR imaging

1.5.1.2 Single exposure

This works thanks to the fact that in CMOS sensors, a bit of computation can be done on the pixels themselves. Each pixel can be read multiple times during the exposure time (each one has different “filling” times), and the values can be combined to obtain a single value for that pixel. This allows to obtain a single image with a high dynamic range. A problem rises once the image must be read from the sensor, we could have more bits per pixel than the sensor can read. We can linearly map the values into the range of the sensor, but this might be hard to do with a really small amount of transistors per pixel. Typically, the values are mapped into a linear-logarithmic scale, with can be easily implemented in hardware on the pixels.

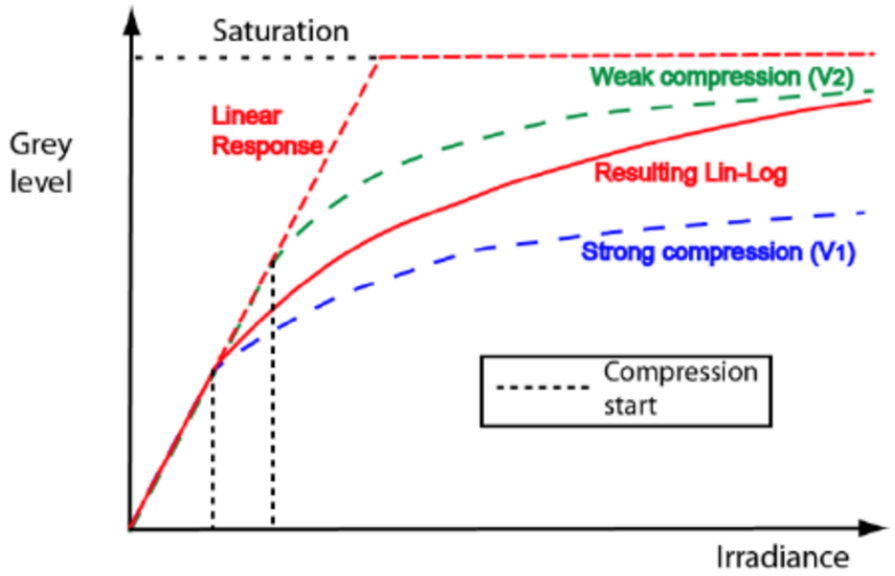


Figure 19: Single exposure HDR imaging

Please note that this introduces a non-linear mapping of the values, which is a change of semantics of the values, and therefore a change of the meaning of the image.

1.5.2 Color cameras

There are different approaches to color cameras:

1.5.2.1 Color filter array

Typically different color channels are captured by different sensors, and the final image is obtained by combining the three channels. This is done by using a color filter array (CFA) that allows each sensor to capture only one color channel (red, green, or blue). The most common CFA is the Bayer filter, which has a pattern of red, green, and blue filters arranged in a specific way. The raw data from the sensors is then processed using a demosaicing algorithm to reconstruct the full-color image.

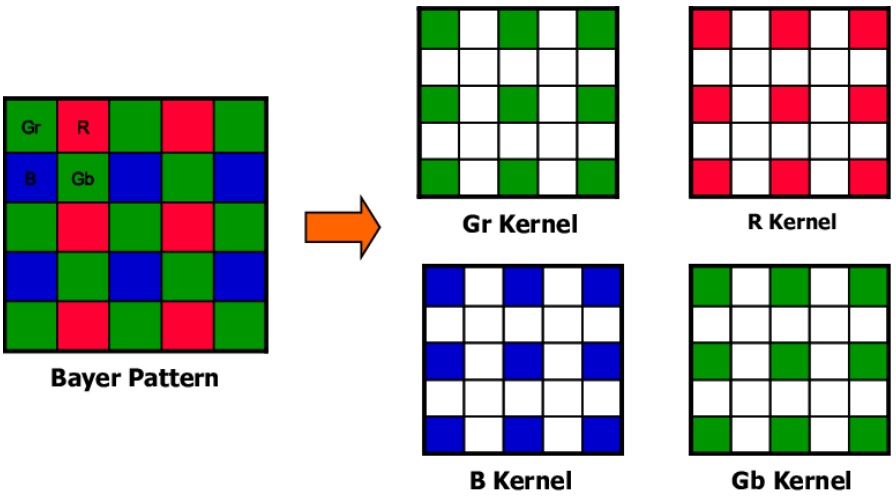


Figure 20: Bayer filter pattern

Quantum efficiency is not defined in the same way for color cameras, since it depends on the wavelength of the light. The colored image is obtained by interpolating the values of the three channels, which are not sampled at the same points. This process is called **demosaicing**. There are many algorithms for demosaicing. The choice of algorithm can affect the quality of the final image, particularly in terms of color accuracy and the presence of artifacts. In general, all demosaicing algorithms introduce artifacts to some level.

Three sensors are sufficient because the color space is three-dimensional, and red, green and blue are a base of the color space.

Remark 1.9. So in practice, when creating a color image, we are sub-sampling the color channels, and then interpolating them to obtain a full-color image.

1.5.2.2 3-CCD cameras

In 3-CCD cameras, the light from the scene is split into three separate beams using a prism or dichroic mirrors, and each beam is directed to a separate CCD sensor that captures one of the primary colors (red, green, or blue). This allows for higher color fidelity and resolution compared to single-sensor cameras with color filter arrays. The three separate images are then combined to create the final color image. 3-CCD cameras are often used in professional video and photography applications. This method avoids the need for demosaicing, as each color channel is captured at full resolution, resulting in better color accuracy and reduced artifacts. But the sensor must be kept in the same position with a micrometer precision, and the prism must be aligned with the sensors. This makes the camera more expensive and less robust to shocks. Also, some of the energy is lost in the prism, and the camera is bigger and heavier than a single-sensor camera.

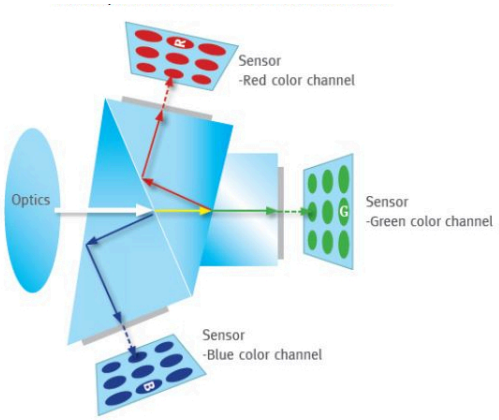


Figure 21: 3-CCD camera

1.5.2.3 Stacked sensors (Foveon)

Three different layers of sensors are stacked on top of each other, with each layer sensitive to a different range of wavelengths (red, green, or blue). This allows for full-color capture at each pixel location without the need for demosaicing. Stacked sensors can provide high color fidelity and resolution, but they can be more complex and expensive to manufacture compared to traditional single-layer sensors with color filter arrays.

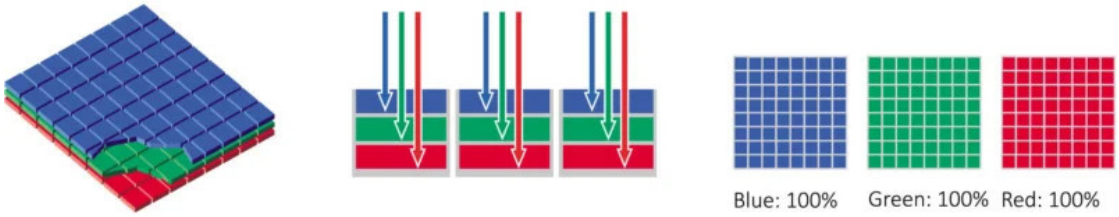


Figure 22: Foveon sensor

1.6 3D cameras

A 3D camera outputs a 4-channel image, with the first three channels being the color channels, and the fourth channel being the depth channel (D). The depth channel contains the distance

of each pixel from the camera. This allows to obtain a 3D representation of the scene. Alternatively, in addition to the color channels, the camera can output a 3D point cloud, which is a set of points in 3D space that represent the scene. Each point has a position in 3D space and a color value. This allows to obtain a 3D representation of the scene.

1.6.1 Time-of-flight cameras

We will focus on time-of-flight (ToF) cameras, which are a type of 3D camera that measures the distance to objects in the scene by calculating the time it takes for a signal to travel from the camera to the object and back. ToF cameras **emit a signal** and measure the time it takes for the signal to return after reflecting off objects in the scene. This time measurement is then used to calculate the distance to each point in the scene, creating a depth map.

> Light is fast, too fast...

Doing a TOF camera with light is possible, but it is very expensive (LIDAR). Some other devices work thanks to amplitude modulation, where the emitter emits a modulated signal, and the sensor measures the phase shift of the returning signal. This phase shift is proportional to the distance to the object, allowing for depth calculation. This type of calculation is really cheap and can be done by standard hardware. This approach also gives me control over the frequency of the signal that is emitted.

1.7 Defects in images

Here we introduce some defects that can be present in images. Defects can be caused by the sensor, the lens, or the environment.

1.7.1 Defects caused by the sensor

1.7.1.1 Blooming

Blooming occurs when the light source illuminating a collection site is so bright that the charge stored at that site overflows into adjacent ones.



Figure 23: CCD blooming and smear. (a) A CCD captured image with blooming and smear, (b) another scene image with blooming and smear.

1.7.1.2 Smear

Smear denotes the undesired bright sections above and below the bright light source, caused by charges' accumulation of the light source during the vertical transfer process.

1.7.2 Defects caused by the lens

The field of view of a camera is the portion of scene space that actually projects onto the retina of the camera. It is not defined by the focal length alone, but also depends on the effective area of the retina (e.g., the area of film that can be exposed in a photographic camera, or the area of the CCD sensor in a digital camera.

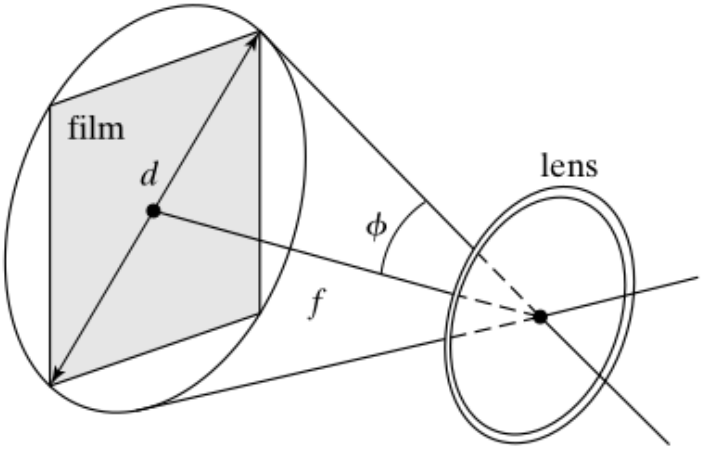


Figure 24: Field of view of a camera

Lenses can introduce distortions in the image, which generally increase with the distance from the optical axis. Therefore, it's better to put the sensor well inside the circle in which the lens projects the image. The most common distortions are:

- **Barrel distortion**: the image magnification decreases with distance from the optical axis, causing straight lines to appear curved outward. This is common in wide-angle lenses.
- **Pincushion distortion**: the image magnification increases with distance from the optical axis, causing straight lines to appear curved inward. This is common in telephoto lenses.

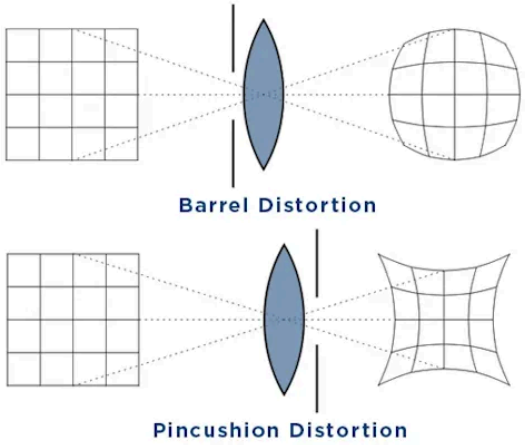


Figure 25: Lens distortion

We can also discriminate between radial and tangential distortion. Radial distortion is caused by the shape of the lens, while tangential distortion is caused by the misalignment of the lens elements.



Figure 26: Radial and tangential distortion

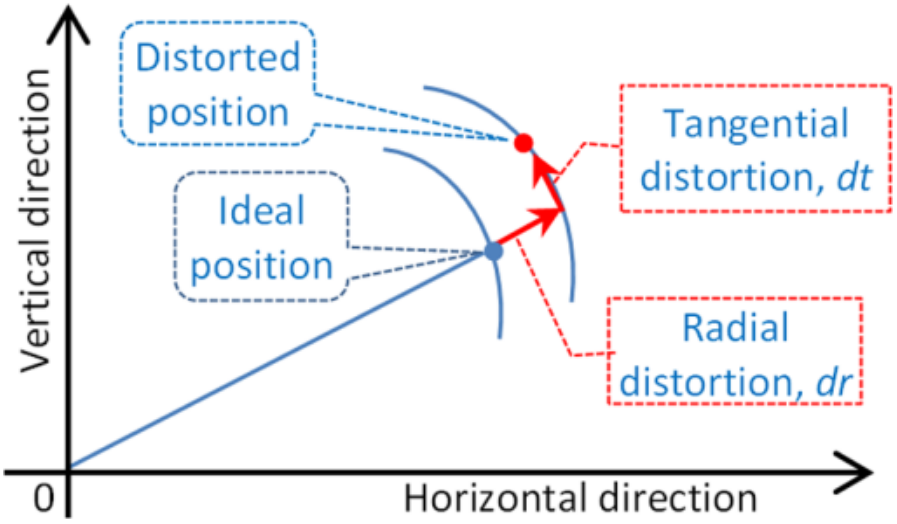


Figure 27: Definitions of radial and tangential distortion on the 2-D plane

Note that both radial and tangential are a **function of the radius**, that is the distance from the optical axis (principal point). Therefore, they are both zero at the optical axis and increase with distance from it.

1.7.2.1 Chromatic aberration

Chromatic aberration is a type of distortion in which there is a failure of a lens to focus all colors to the same convergence point. It occurs because lenses have different refractive indices for different wavelengths of light, causing them to bend light of different colors by different amounts. This results in a misalignment of colors in the image, leading to color fringing around the edges of objects, especially in high-contrast areas. Chromatic aberration can be minimized by using achromatic lenses, which are designed to bring two wavelengths (typically red and blue) into focus in the same plane, or by using apochromatic lenses, which bring three wavelengths into focus in the same plane.

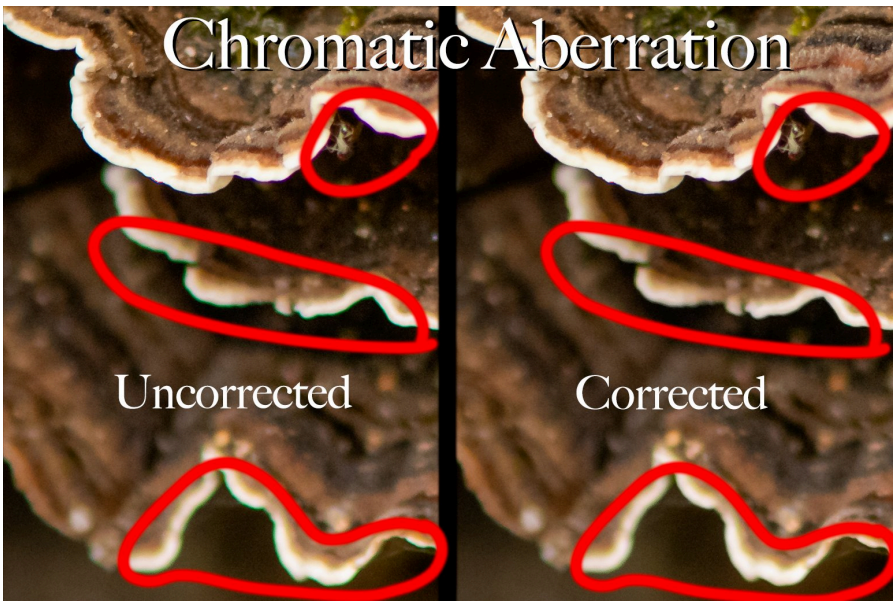


Figure 28: Chromatic aberration example

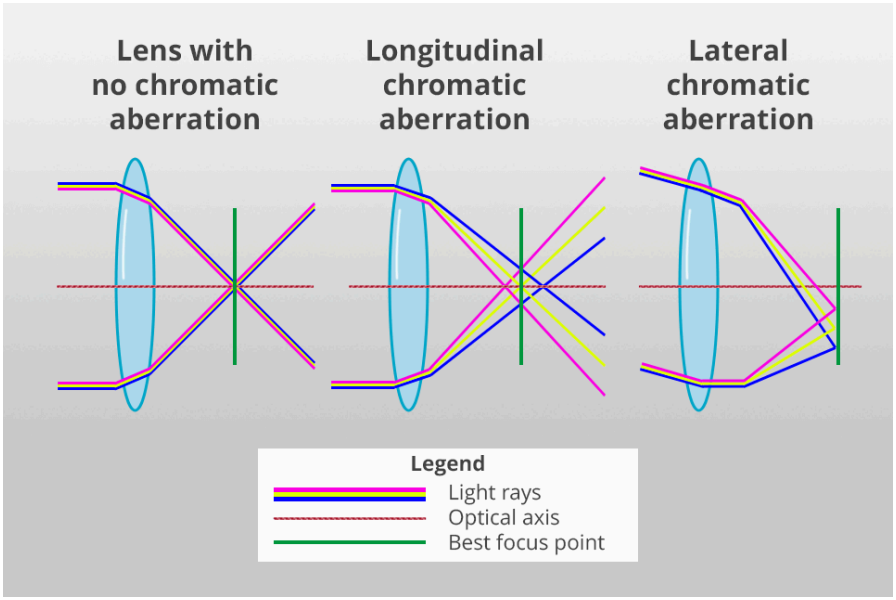


Figure 29: Chromatic aberration types

1.7.2.2 Vignetting

Light beams emanating from object points located off-axis are partially blocked by the various apertures (including the individual lens components) positioned inside the lens to limit aberrations. This phenomenon, called **vignetting**, causes the brightness to drop in the image periphery. Vignetting may pose problems to automated image analysis programs, but it is not as important in photography thanks to the human eye's remarkable insensitivity to smooth brightness gradients.

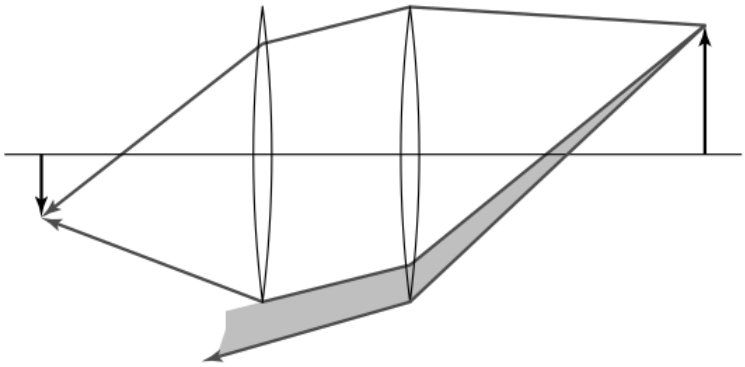


Figure 30: Vignetting effect in a two-lens system. The shaded part of the beam never reaches the second lens. Additional apertures and stops in a lens further contribute to vignetting.



Figure 31: Example of vignetting in a photograph. The corners of the image are darker than the center.

B. Mathematical concepts

Here we present some related mathematical concepts that are required for the course.

B.1. Roto-translation

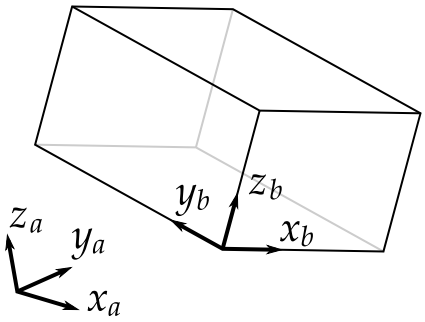


Figure 32: Rigid body attached frame. The position and orientation of a rigid body can be described by attaching a frame to it.

Consider a rigid body as of figure Figure 32, rigid means that the points of the body do not change their relative position. The position and orientation of a rigid body can be described by attaching a frame to it. The coordinates of the points of the object are expressed in the attached frame (x_b, y_b, z_b) and we would want to express them in the reference frame (x_a, y_a, z_a).

In many scientific fields, the **degrees of freedom** of a system is the number of parameters of the system that may vary independently. For example, a point in the 3D space has three degrees of freedom for translation: its coordinates; In 3D space, n points have $3n$ degrees of freedom for translation. If the object is rigid, that is not the case, because the points of the object cannot move independently. In fact, a rigid body has 6 degrees of freedom: 3 for translation and 3 for rotation.

The position and orientation of a rigid body can be described by a roto-translation matrix T that transforms the coordinates of the points of the object expressed in the attached frame (x_b, y_b, z_b) to the reference frame (x_a, y_a, z_a): The union of the position and the orientation of a rigid body is called the **pose** of the rigid body.

B.2. Euler angles

Euler angles describe the orientation of a rigid body in space through three successive intrinsic rotations:

- **Precession (φ):** Rotation around the initial z-axis by an angle φ. The system (x, y, z) becomes (x', y', z'), with $z' = z$.
- **Nutation (θ):** Rotation around the intermediate x'-axis by an angle θ. The system becomes (x'', y'', z''), with $x'' = x'$.
- **Proper rotation (ψ):** Rotation around the intermediate z''-axis by an angle ψ. The final system is (x''', y''', z'''), with $z''' = z''$.

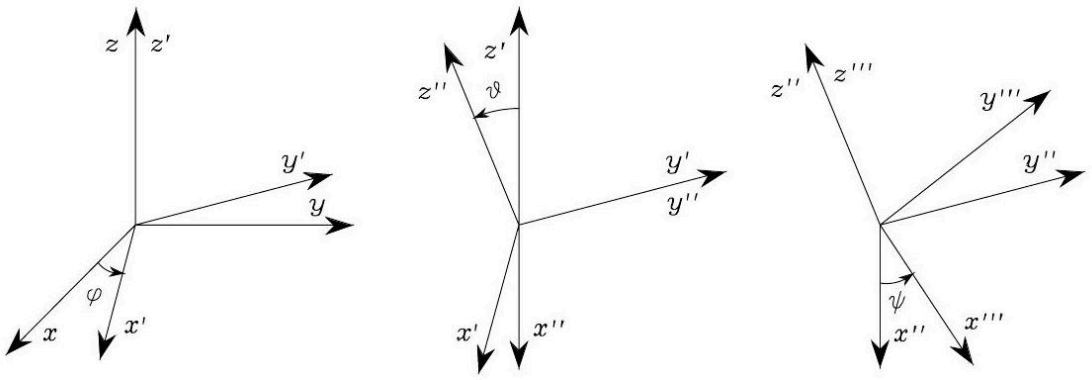


Figure 33: Euler angles. The orientation of a rigid body in space can be described through three successive intrinsic rotations: precession, nutation and proper rotation.

B.3. Roll, pitch and yaw angles

Roll, pitch and yaw angles describe the orientation of a rigid body in space through three successive extrinsic rotations: Roll, pitch and yaw.

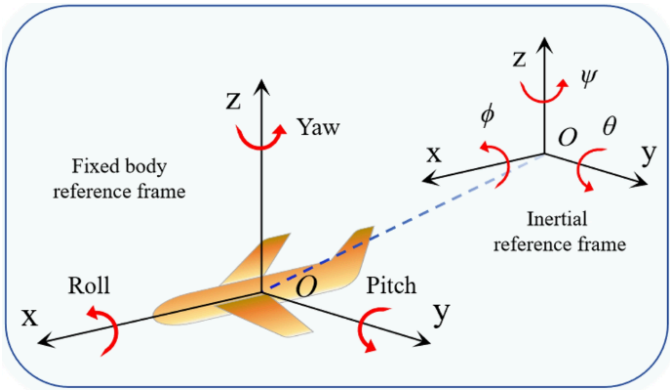


Figure 34: Roll, pitch and yaw angles.

B.4. Roto-translation matrix

Let's combine the translation and rotation of a rigid body in a single transformation matrix T called roto-translation matrix, which represents the relative pose between the relative frame of the object and the reference frame.

Take for example a frame $O\text{-}zyx$, we rotate it around the z-axis by an angle α, we can say that α is the value of one of the degrees of the change of orientation between the two frames.

$$x' = \begin{pmatrix} \cos\alpha \\ \sin\alpha \\ 0 \end{pmatrix} \quad y' = \begin{pmatrix} -\sin\alpha \\ \cos\alpha \\ 0 \end{pmatrix} \quad z' = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} \tag{20}$$

Then, the rotation matrices are defined as:

$$\begin{aligned} R_z(\alpha) &= \begin{pmatrix} \cos\alpha & -\sin\alpha & 0 \\ \sin\alpha & \cos\alpha & 0 \\ 0 & 0 & 1 \end{pmatrix} \\ R_y(\beta) &= \begin{pmatrix} \cos\beta & 0 & \sin\beta \\ 0 & 1 & 0 \\ -\sin\beta & 0 & \cos\beta \end{pmatrix} \\ R_x(\gamma) &= \begin{pmatrix} 1 & 0 & 0 \\ 0 & \cos\gamma & -\sin\gamma \\ 0 & \sin\gamma & \cos\gamma \end{pmatrix} \end{aligned} \tag{21}$$

We postmultiply the coordinates of a point expressed in the frame $O\text{-}zyx$ by the rotation matrix to obtain the coordinates of the point expressed in the rotated frame $O\text{-}z'y'x'$. If we compose the rotations, we can compute one after the other, for example, we can first rotate around the z-axis and then around the y-axis:

$$\begin{pmatrix} x'' \\ y'' \\ z'' \end{pmatrix} = R_y(\beta) R_z(\alpha) \begin{pmatrix} x \\ y \\ z \end{pmatrix} \tag{22}$$

We can represent translation and rotation in a single transformation matrix T called roto-translation matrix, which represents the relative pose between the relative frame of the object and the reference frame. To do this, we can use homogeneous coordinates, which allow us to represent both rotation and translation in a single matrix. So given a point $p = \begin{pmatrix} x \\ y \\ z \end{pmatrix}$ we have $\tilde{p} = \begin{pmatrix} w_1 \\ w_2 \\ w_3 \\ w_4 \end{pmatrix}$ with

$$w_1 = \frac{x}{w_4}, w_2 = \frac{y}{w_4}, w_3 = \frac{z}{w_4} \tag{23}$$

and we can express the roto-translation as:

$$\tilde{p}^0 = A_1^0 \tilde{p}^1 \tag{24}$$

with

$$A_1^0 = \begin{pmatrix} R_1^0 & o_1^0 \\ 0^T & 1 \end{pmatrix} \tag{25}$$

where R_1^0 is the rotation matrix from frame 1 to frame 0, and o_1^0 is the translation vector of the origin of frame 1 with respect to frame 0. The rotation gives three degrees of freedom, while

the translation gives three more degrees of freedom, for a total of six degrees of freedom for a rigid body.